

Writing a reflection engine from scratch

Manu Sanchez @Manu343726

Reflection

From wikipedia

In computer science, reflection is the ability of a computer program to examine, introspect, and modify its own structure and behavior at runtime

Reflection

- Examine the program structure
- Manipulate the program structure programmatically
- At compile time: "***static reflection***"
- At runtime: "***dynamic reflection***"

Given...

```
namespace mylib
{
    class MyClass
    {
    public:
        MyClass() = default;

        void myMethod(const std::string& str);

        int i = 0;
    };
}
```

Static reflection

```
using Method = cpp::static_reflection::Method<
    decltype(&MyClass::myMethod),
    &MyClass::myMethod
>;

// File where the method is declared
constexpr auto filePath = Method::SourceInfo::file();
```

- Pros:
 - Type safe
 - Zero overhead
- Cons:
 - Ugly for the user. *TMP based* (`constexpr` could help)

Reminder...

```
namespace mylib
{
    class MyClass
    {
    public:
        MyClass() = default;

        void myMethod(const std::string& str);

        int i = 0;
    };
}
```

Dynamic reflection

```
cpp::dynamic_reflection::Runtime runtime;
...
const auto& myObject = runtime
    .namespace_()           // Global namespace
    .namespace_("mylib")    // ::myLib namespace
    .class_("MyClass")      // ::myLib::MyClass class
    .create();              // Returns object

// Invoke method on the object:
myObject("myMethod")(std::string{"hello"});

// Set member variable:
myObject["i"] = 42;
```

- Pros:
 - User-friendlier API
 - Entities manipulable at runtime
- Cons:

Our goal

Giving both static and dynamic reflection features to C++

No, really

- Let's restrict the API to OOP:
 - Classes
 - Enums
 - Public non static member functions (*methods*)
 - Public non static member variables (*fields*)

What we need

- Entity metadata:
 - Classes declared in a header
 - Class members
 - Source file path
 - Source line
 - Pointers to members
- An API to read metadata: `api(entity) -> metadata`
 - A C++ representation of the metadata
 - Entity to metadata mapping

Entity: A C++ sourcecode semantic element (A class, member function, enum, namespace, etc)

Collecting metadata

Collecting metadata

- Per-header metadata

For each header file generate a file with the metadata collected from the C++ header

- No global codebase information

No cross-header namespace info, but makes things easier

- Metadata is stored in **generated C++ header files**

Easy to import, just `#include`

- *So metadata is stored as C++ code? Exactly. More on this later*

Collecting metadata: libclang

libclang

- A C interface to the clang library
- Provides support for C++ AST parsing and transversal

libclang: How it works?

- The Clang compiler is written as a series of C++ libraries
 - **-pedantic** **note:** Clang is not a compiler but a C++ frontend for LLVM
- libclang provides a C interface for the clang API
- The libclang C API is **stable between clang releases**
 - *The C++ API may apply breaking changes between versions*
- C ABI: **There are libclang bindings for multiple languages**

libclang: Hello world with Python

```
from clang.cindex import Index, Cursor, CursorKind

args = ['-std=c++11'] # Compilation options
index = Index.create() # Create a compilation DB

# Parse a translation unit:
tu = index.parse('project.hpp', args)

def visit(cursor):
    print 'file {}:{} {} ({}).format(
        cursor.location.file,
        cursor.location.line,
        cursor.spelling,
        cursor.kind)

    for c in cursor.children():
        visit(c)

visit(tu) # Visit the translation unit AST
```


libclang python basics

Python wraps the AST visitation API as a `Cursor` class with properties about the current AST node being visited:

- `kind` : Type of node (A function declaration, a class declaration, a namespace declaration, etc). `CursorKind` enum.
- `spelling` : Name of the entity being declared/defined
- `displayname` : Detailed entity name. `
- `location` : Entity location in the sourcecode: File, line, etc

libclang python basics

Many more properties. See: <https://github.com/llvm-mirror/clang/blob/master/bindings/python/clang/cindex.py>

(Back to) Reading metadata

Reading metadata

- Take only the interesting cursors (those of an specific kind):
`CursorKind.CLASS_DECL` , `CursorKind.STRUCT_DECL` ,
`CursorKind.ENUM_DECL` , `CursorKind.CXX_METHOD` ,
`CursorKind.FIELD_DECL` , etc
- Expose new properties you might need:

```
@property
def fullname(self):
    if self.parent is None:
        return '::' + self.cursor.spelling
    else:
        return self.parent.fullname + '::' +
               self.cursor.spelling
```

Generating C++ code

- Write templates with the C++ model of the metadata
- Pass the processed AST to your favorite template engine
- Done!

Generating C++ code

- The format of your generated metadata depends on the reflection API you are going to write
 - See Qt's MOC compiler, `QMetaType` , and `xxxx_moc.cpp` generated code.
- We are writing an static reflection API, so metadata must be manipulable at compile time
- Which means...

Template metaprogramming!!!

- No, really?
- Well, that's what I did
- Try `constexpr` + `boost::hana`

Generating C++ code: Metadata format

- `cpp::static_reflection::meta` : Templates storing metadata

```
namespace cpp::static_reflection::meta {  
    template<  
        typename Spelling,  
        typename DisplayName,  
        typename FullName,  
        typename FilePath,  
        std::size_t Line  
        ...  
    > class SourceInfo { ... };  
}
```

Generating C++ code: Metadata format

- Public static constexpr member functions to return the properties:

```
template<...>
class SourceInfo
{
public:
    // Returns a const ref to the constexpr string
    // with the full name:
    static constexpr decltype(auto) fullName()
    {
        return cpp::constexpr::SequenceToString<
            FullName
        >::str();
    }
    ...
};
```


Generating C++ code: Metadata format

- Public typedefs to typelists to represent sets of metadata:

```
template<
    // A SourceInfo<...> instance:
    typename SourceInfo_,
    // A typelist of Function<fptr type, fptr>
    // instances:
    typename Methods_,
    ...
> class Class
{
public:
    using SourceInfo = SourceInfo_;
    using Methods = Methods_;
    ...
};
```

- `cpp::static_reflection::meta` templates have another metadata instances (or lists of) as parameters

Generating C++ code: Writing metadata

- Write a template that takes a processed AST and writes `cpp::static_reflection::meta` instances into a header file.

Remember when we said "Per-header info only?"

Generating C++ code: Writing metadata

```
{% macro sourceinfo(node, kind, class = None) -%}  
::cpp::static_reflection::meta::SourceInfo<  
    {{sourceinfo_entity(node, class)}},  
    ::cpp::static_reflection::Kind::{{kind}},  
    ::cpp::meta::string<{{node.fullname_as_charpack}}>,  
    ::cpp::meta::string<{{node.spelling_as_charpack}}>,  
    ::cpp::meta::string<{{node.displayname_as_charpack}}>  
    ::cpp::meta::string<{{node.file_as_charpack}}>,  
    {{node.cursor.location.line}}  
>  
{%- endmacro %}
```

See [siplasplas@github/src/reflection/parser/templates/](https://github.com/siplasplas/src/reflection/parser/templates/)

The static reflection API

The static reflection API: Goals (again)

- Mapping from entity to metadata
- Metadata can be queried at compile-time

The static reflection API: API entry points

- How we access `cpp::static_reflection::meta` instances?
- `cpp::static_reflection namespace: Class , Enum , Field ...` All counterparts of `meta` templates that take **an entity as parameter**:

```
#include <path/to/gen/code/myclass.hpp>

using MyClassData = cpp::static_reflection::Class<
    MyClass
>;
```

The static reflection API: API entry points

- When generating code, we specialize `cpp::static_reflection::codegen` templates with the entities as parameter.
- Those specializations inherit from the metadata (`cpp::static_reflection::meta` instances)
- `cpp::static_reflection::codegen` instead of `cpp::static_reflection` to provide an extra layer of indirection
 - So we can get rid of all the libclang stuff when standard reflection becomes real

The static reflection API: API entry points

```
namespace cpp::static_reflection::codegen {  
    template<>  
        class Field<  
            decltype(&::project::Dog::dogAlias),  
            &::project::Dog::dogAlias  
        > :  
            public ::cpp::static_reflection::meta::Field<  
                ::cpp::static_reflection::meta::SourceInfo<  
                    ::cpp::static_reflection::codegen::Field<  
                        decltype(&::project::Dog::dogAlias),  
                        &::project::Dog::dogAlias  
                    >,  
                    ::cpp::static_reflection::Kind::FIELD,  
                    ::cpp::meta::string<'::', ':', 'p', 'r', '  
                    ::cpp::meta::string<'d', 'o', 'g', 'A', '  
                    ::cpp::meta::string<'d', 'o', 'g', 'A', '  
                    ::cpp::meta::string<('/', 'h', 'o', 'm', '  
                47  
                >, ...  
            >  
        {} } // C++17 only...
```


Playing with the static reflection API: JSON serialization

JSON Serialization

- Let's write two template functions:

- `serialize(const T& value)`

- `deserialize<T>(const nlohmann::json& json)`

This one must be implemented through a template class,
(we cannot apply SFINAE directly)

- Completely generic
- Type safe
- JSON for Modern C++

JSON Serialization

- Map C++ values to JSON objects `{"type": "typename", "value": ...}`
- The type of the original value is stored as a string. We can check it later

A type name string can be "easily" computed at compile time. See <https://github.com/Manu343726/ctti>
- Each kind of C++ value is serialized in an optimal way

JSON Serialization: Fundamental types

```
template<typename T>
std::enable_if_t<std::is_fundamental<T>::value, nlohmann::
serialize(const T& value)
{
    auto json = nlohmann::json::object();

    json["type"] = cpp::lexical_cast(ctti::type_id(value)
    json["value"] = value;

    return json;
}
```

JSON Serialization: Fundamental types

```
template<typename T>
class Deserialize<T, std::enable_if_t<std::is_fundamental
{
public:
    static T apply(const nlohmann::json& json)
    {
        SIPLASPLAS_ASSERT_EQ(
            json["type"].get<std::string>(),
            cpp::lexical_cast(ctti::type_id<T>().name())
        );

        return json["value"];
    }
};
```

JSON Serialization: unordered_map

```
template<typename Key, typename Value>
nlohmann::json serialize(const std::unordered_map<Key, Value>& value)
{
    auto json = nlohmann::json::object();
    auto array = nlohmann::json::array();

    for(const auto& keyValue : value)
    {
        const auto& key = keyValue.first;
        const auto& value = keyValue.second;

        array.push_back(nlohmann::json::object({
            {"key", serialize(key)},
            {"value", serialize(value)}
        }));
    }

    json["type"] = cpp::lexical_cast(ctti::type_id(value));
    json["value"] = array;

    return json;
}
```

JSON Serialization: unordered_map

```
template<typename Key, typename Value>
class Deserialize<std::unordered_map<Key, Value>, void>
{
public:
    static std::unordered_map<Key, Value> apply(const nlohmann::json& json)
    {
        SIPLASPLAS_ASSERT(json["value"].is_array());
        SIPLASPLAS_ASSERT_EQ(
            json["type"].get<std::string>(),
            cpp::lexical_cast<ctti::type_id<
                std::unordered_map<Key, Value>
            >().name())
    );
    }
```

```
std::unordered_map<Key, Value> map;

for(const auto& keyValue : json["value"])
{
    // Weird "I was debugging" checks here...
    SIPLASPLAS_ASSERT(keyValue.is_object())(
        "keyValue ({})) must be an object node. Is
        keyValue.dump(),
        keyValue.type()
    );

    map[Deserialize<Key>::apply(keyValue["key"])] =
        Deserialize<Value>::apply(keyValue["value"]);
}

return map;
};
```


And finally...

JSON Serialization: Objects

```
template<typename T>
std::enable_if_t<std::is_class<T>::value, nlohmann::json>
serialize(const T& object)
{
    auto json = nlohmann::json::object();
    auto fields = nlohmann::json::object();
```

```

// User defined classes are serialized by serializing
// generating a "<member name>: <value>" JSON object
// serialized values
cpp::foreach_type<
    typename cpp::static_reflection::Class<T>::Fields
>([&](auto type)
{
    using FieldInfo = cpp::meta::type_t<decltype(type)

    // We use the spelling of the field as key, and t
    fields[cpp::lexical_cast(FieldInfo::SourceInfo::s
        cpp::invoke(FieldInfo::get(), object) // C++1

    );
});

json["type"] = cpp::lexical_cast(ctti::type_id(object
json["value"] = fields;

return json;
}

```

JSON Serialization: Objects

```
template<typename T>
class Deserialize<
    T,
    std::enable_if_t<
        std::is_class<T>::value &&
        IsNotSpecializedType<T>::value // std::string,
                                         // std::tuple,
                                         // etc
    >
>
{
public:
```

```
static T apply(const nlohmann::json& json)
{
    SIPLASPLAS_ASSERT_EQ(
        json["type"].get<std::string>(),
        cpp::lexical_cast(ctti::type_id<T>().name())
    );

    const auto& fields = json["value"];
    T object;
```

```

cpp::foreach_type<
    typename cpp::static_reflection::Class<T>::Fields
>([&](auto type)
{
    using FieldInfo = cpp::meta::type_t<decltype(type)
    using Type = typename FieldInfo::value_type;

    cpp::invoke(FieldInfo::get(), object) =
        Deserialize<Type>::apply(fields[cpp::lexical_
            FieldInfo::SourceInfo::spelling()
        ]));
});

return object;
}
};

```

JSON Serialization

See the full example at:

```
siplasplas@github/examples/reflection/static/serializatio  
n.cpp
```

Google Protobuf comm model to C++ data model

- Given a communication model written using Protobuf, and a C++ business data model
- **Can we map between these two models automatically?**

Google Protobuf comm model to C++ data model

- Why?
- [WM&IP](#) aka "*telefonillo*". We have a 2k LOC file **just for this**.
- I'm sick of doing this stuff again and again...

Google Protobuf comm model to C++ data model

- We invoke two `serialize<Protobuf, T>(const T& value)` and `deserialize<T, Protobuf>(const Protobuf& proto)`
- Very similar to the JSON example
- Users can define conversions through a `cast<From, To>(const From& from)` customization point

Google Protobuf comm model to C++ data model: Algorithm

- To serialize (C++ to protobuf):
 - i. Foreach field in the C++ object...
 - ii. Search for matching setters in the protobuf c++ generated code: `set_xxxx()` , `set_allocated_xxxx()` , etc
 - iii. Get the nearest setter. Invoke it and pass the value of the field. Apply conversion if needed
- To deserialize (Protobuf to C++):
 - i. Default construct a C++ object
 - ii. Foreach field in the C++ object...
 - iii. Search for matching getters in the protobug c++ generated code: `xxxx()` , `get_xxxx()` , etc

Google Protobuf comm model to C++ data model

Full example at

`siplasplas@github/examples/reflection/static/protoserialization`

Dynamic reflection

Dynamic reflection

- Parsing and accessing metadata at compile time is the hard thing
- Dynamic reflection just exposes this information at runtime

Dynamic reflection

- `cpp::dynamic_reflection` exposes reflection information in the form of a classic inheritance-based entity system
 - `Class` , `Field` , `Enum` ; inheriting from `Entity`
 - Entities are managed by a `Runtime`

Dynamic reflection: Entities

```
namespace cpp::dynamic_reflection {  
  
class Entity : public std::enable_shared_from_this<Entity  
{  
public:  
    virtual ~Entity() = default;  
  
    Entity& parent();  
    Runtime& runtime();  
  
    const SourceInfo& sourceInfo() const;  
    const std::string& name() const;  
    const std::string& fullName() const;  
    const SourceInfo::Kind& kind() const;  
  
    bool detached() const;  
    void attach(Runtime& runtime);  
    bool orphan() const;  
};
```



```

void attach(const std::weak_ptr<Entity>& parent);
void addChild(const std::shared_ptr<Entity>& entity);

bool isChild(const std::shared_ptr<Entity>& entity) const;
bool isChild(const Entity& entity) const;
bool isChildByName(const std::string& name) const;
bool isChildByFullName(const std::string& fullName) const;
Entity& getChildByFullName(const std::string& fullName);
Entity& getChildByName(const std::string& name);
std::vector<std::string> getChildrenNamesByKind(const std::string& kind);

std::shared_ptr<Entity> pointer();

friend bool operator==(const Entity& lhs, const Entity& rhs);
friend bool operator!=(const Entity& lhs, const Entity& rhs);

protected:
    Entity(const SourceInfo& sourceInfo);

private:
    ...
};

```

Dynamic reflection: Entities

- All entities are registered **by full qualified name** in their runtime
- Entities maintain a hierarchical relationship, so users can walk along the namespaces, classes, etc
- Entities can be queried by simple name (spelling) **through the parent entity**
- The global namespace is the root of the hierarchy

Dynamic reflection: Entities

- There's an `Entity` subclass for each kind of entity supported (`Class`, `Enum`, `Namespace`, etc)
- Entity classes provide factories to return entity objects from type-erased entity pointers:

```
// Throws if there's no "::mylib::MyClass" entity or  
// a class entity  
Class& class_ = Class::fromEntity(runtime.get("::myli
```

- Entity classes provide operations specific to that kind of entity:

```
auto object = class_.create(); // Default constructs  
Enum& enum_ = class_.enum_("Enum"); // Return ::mylib  
  
std::int64_t firstEnumVal = enum_.value(0);  
const std::string& enumValStr = enum_.toString(firstE  
std::int64_t enumValue = enum_.fromString("FIRST");
```

Dynamic reflection: Objects

- The `Class` entity returns instances of the class
- The instances are stored in a type-erased container,
`cpp::Any32`

Dynamic reflection: Objects

- `cpp::Any32` :
 - Can store values of different types at runtime
 - Similar to C++17 `std::any` or `boost::any`
 - 32 byte fixed-size storage with fallback to dynamic allocation
 - Stores a map of **attributes and functions** for easy OO-like manipulation of the object

Dynamic reflection: Objects

```
cpp::Any32 object = class_.create();  
  
// Invoke object.method with params 42 and "hello"  
object("method")(42, std::string{"hello"});  
  
// Assign a value to the public member variable "i"  
object["i"] = 42 * 42;
```

Dynamic reflection: Runtime API loading

- Dynamic libraries can export their API by invoking
`SIPLASPLAS_REFLECTION_DYNAMIC_EXPORT_TYPES(Ts...)`
macro
- The macro uses static reflection to collect reflection information of the specified library types
- That information is accesible through a C API function that loads it into a dynamic reflection

`Runtime` object

```
// mylib_export.cpp

#include <reflection/mylib.hpp>
#include <siplasplas/reflection/dynamic/exporttypes.hpp>

SIPLASPLAS_REFLECTION_DYNAMIC_EXPORT_TYPES(
    ::mylib::MyClass,
```

Dynamic reflection: Runtime API loading

- Users can load APIs at runtime with a `cpp::dynamic_reflection::RuntimeLoader` object, which takes a dynamic library handle and returns a `Runtime` object

Dynamic reflection: Runtime API loading

```
#include <siplasplas/reflection/dynamic/runtimeloader.hpp>

int main()
{
    auto lib = cpp::DynamicLibrary::load("libmylib.so");
    cpp::dynamic_reflection::RuntimeLoader loader{lib};
    auto& runtime = loader.runtime();

    auto object = runtime
        .namespace_()
        .namespace_("mylib")
        .class_("MyClass")
        .create();

    object("run")(); // Invoke object.run()
}
```

Future work

Future work: A dynamic type system

- Dynamic type conversions
- Declaring "metatypes" at runtime

Future work: Runtime C++ compilation

- Proof of concept at:

`siplasplas@github/examples/cmake/project.cpp`

- See `siplasplas-cmake` module in the docs
 - Really looking forward to put my hands on cmake-server!

Future work: Maintenance

- I've had no time to handle some things properly
- From June to middle July I was full(siplasplas)-time writing docs :(
- This used to be fully cross-platform (VS2015 included) but I broke a lot of things when writing siplasplas-typeerasure...
- The CI builds are broken (*It works on my machine*)
- Third party management
 - Github says I have 5% of cmake scripts in the project...
 - conan.io save me. Diego you owe me a master class!
- Devlog and more crazy stuff blogposts!

- <https://github.com/Manu343726/siplasplas>
- Website
- Latest doxygen docs

thanks_for_comming(); }